

Assignment 1

FIT5225 2025 T3

CloudPose:

*Creating and Deploying a Pose Detection/Estimation
Web Service within a Containerised Environment in Clouds*

Due date: Tuesday, 5 August at 11:55 PM

1 Synopsis and Background

Pose estimation is a computer vision technique for identifying the location and orientation of an object in an image or video. It typically involves detecting specific key points on the object, such as joints for humans or landmarks for other objects. Pose estimation typically uses machine learning models, often deep neural networks, to learn the patterns and relationships between pixels in an image and the location of key points. These models are trained on large datasets of images with annotated key points.

This project aims to build a web-based system hosted in Oracle Cloud Infrastructure (OCI) that we call **CloudPose**. It allows end-users to send an image to a web service hosted by containers and receive a list of key points detected and an annotated image in their request.

The project will perform the required operations using several pre-trained pose estimation models and libraries. The model and your web service will be hosted as containers in a Kubernetes cluster as the container orchestration platform. The web service is also designed to be a RESTful API that utilises Python's FASTAPI library. We are interested in examining CloudPose's performance by varying the rate of requests sent to the system (demand) using load generation tools like Locust and the number of existing Pods within the Kubernetes cluster (resources).

This assignment has the following objectives:

- Writing two Python **web services** that accepts images in JSON object format, uses the provided models and libraries to process images, and returns 1) a JSON object with a list of detected key points; 2) An annotated image of detected keypoints.
- Building a **Docker Image** for the web service.
- Creating a **Kubernetes cluster** on virtual machines (instances) in Oracle Cloud Infrastructure (OCI). Ali cloud can be as a substitute (but you should illustrate how. The Oracle used in the following tasks can be all replaced by Ali cloud if you want).
- Deploying a **Kubernetes service** to distribute inbound requests among pods that are running the pose estimation service.
- Writing a **load generation** scripts using Locust.
- **Testing** the system under varying load and number of pods conditions.

You can focus on these objectives one after another to secure partial marks.

The models and libraries: We provide four different pre-trained models and corresponding example codes to use them. **Based on the last digit of your student number, download one of the following models to build your web service.**

- If your student number ends with 0 or 5: Use Model 1
- If your student number ends with 1, 4, 6: Use Model 2
- If your student number ends with 2 or 7: Use Model 3
- If your student number ends with 3, 8 or 9: Use Model 4

2 The web service - [20 Marks]

2.1 Pose Estimation JSON API

You are required to develop a RESTful API that allows clients to upload images to the service. You must use **FASTAPI** to build your web service and use any port between **60000-61000**. Your FASTAPI server should be able to handle multiple clients concurrently. Each image should be sent to the web service using an HTTP POST request containing a JSON object with a unique ID (e.g. UUID) and a base64-encoded image. Since an image is binary data, it cannot be directly inserted into JSON. You must convert the image into a textual representation that can then be used as a normal string. The most common way to encode an image into text is using the base64 method. A sample JSON request used to send an image could be as follows:

```
{
  "id": "06e8b9e0-8d2e-11eb-8dcd-0242ac130003",
  "image": "YWRzZmFzZGZhc2RmYXNkZmFzZGYzNDM1MyA7aztqMjUzJyBqaDJzM2 ..."
}
```

The web service creates **a thread per request** and uses your pose estimation model and library to detect key points in the image. As a suggestion, you can begin with the image encoding part of the JSON message, consider developing your web service and testing it with basic Postman HTTP requests. Once you've confirmed that your web service functions correctly, you can proceed to create your client requests in accordance with the web service API using Locust. For each image (request), your web service returns a JSON object with a list of all key points detected in that image as follows:

```
{
  "id": "The id from the client request",
  "count": n,
  "boxes": [
    {"x": x1, "y": y1, "width": w1, "height": h1, "probability": p1},
    {"x": x2, "y": y2, "width": w2, "height": h2, "probability": p2},
    ...
    {"x": xn, "y": yn, "width": wn, "height": hn, "probability": pn}
  ],
  "keypoints": [
    [[x1, y1, p1],
    [x2, y2, p2],
    ...
  ]
}
```

```

    [x17,y17,p17]],
    ...
],
speed_preprocess: a,
speed_inference: b,
speed_postprocess: c
}

```

- The “id” is the same ID sent by the client along with the image. This is used to associate an asynchronous response with the request on the client side.
- The “count” n represents the number of persons detected in the image/video.
- The “boxes” represent the list of n rectangles around the detected person, with information for (x,y) coordinate, width, height and the probability of the rectangular area containing a person.
- The “keypoints” represent the detected key points in the image/video. All provided models support the detection of at least 17 key points for each person. You don’t need to modify the models, use as many as keypoints provided by your model.

2.2 Pose Estimation Image API

Another web service endpoint is required to provide an annotated response. It takes the same request as the previous API but at a different endpoint, generating an annotated image encoded with base64 as a response. Please see the following sample images:

You are required to use the provided pre-trained models to develop a fast and reliable RESTful API for pose estimation. You will use pre-trained network weights, so there is no need to train the pose estimation program yourself and the required configuration files.

3 Dockerfile - [10 Marks]

Docker builds images by reading the instructions from a file known as *Dockerfile*. Dockerfile is a text file that contains all ordered commands needed to build a given image. You are required to create a Dockerfile that includes all the required instructions to build your Docker image. You can find Dockerfile reference documentation here: <https://docs.docker.com/engine/reference/builder/>.

To reduce complexity, dependencies, file sizes, and build times, avoid installing extra or unnecessary packages just because they might be ‘nice to have.’ For example, you do not need to include a text editor in your image. It is important to optimise your Dockerfile while keeping it easy to read and maintain.

4 Kubernetes Cluster - [10 Marks]

You are tasked with installing and configuring a Kubernetes cluster on OCI VMs. For this purpose, you are going to install K8s on three VM instances on OCI (All your VM instances should have 4GB Memory and 2 OCPUs). Please use the node from HongKong. You need to set up a K8s cluster with 1 controller and 1 worker nodes that run on OCI VMs. You need to install the Docker engine on VMs. You should configure your K8s cluster with Kubeadm.

5 Kubernetes Service - [10 Marks]

After you have a running Kubernetes cluster, you need to create service and deployment configurations that will in turn create and deploy required pods in the cluster. The official documentation of Kubernetes



Figure 1: Original Image



Figure 2: Image with key points annotated

contains various resources on how to create pods from a Docker image, set CPU and/or memory limitations, and the steps required to create a deployment for your pods using selectors. Please make sure to set the CPU request and CPU limit to “0.5” and the memory request and limit to “512MiB” for each pod.

Initially, you will start with a single pod to test your web service and gradually increase the number as described in Section 7. The preferred way of achieving this is by creating replica sets and scaling them accordingly.

Finally, you are required to expose your deployment to enable communication with the web service running inside your pods. You can make use of **Service and NodePort or an Ingress controller** to expose your deployment. You will need to call the pose estimation service from various locations as described in the next section. OCI restricts access to your VMs through its networking security measures. Therefore, you should ensure that your controller instance has all the necessary ports opened and that necessary network configurations, including OCI “Security Lists,” are properly set up. You may also need to open ports on the instance-level firewall (e.g. firewall or iptables).

6 Locust load generation - [10 Marks]

Create a Locust script to simulate concurrent users accessing your RESTful API. Ensure the API can handle the load and respond promptly without crashing or experiencing significant delays. Your next task involves monitoring and recording relevant performance metrics such as response time, query per second (QPS), and error rate during load testing.

First, install Locust (if not already installed) and familiarise yourself with its documentation to create load-testing scenarios. Configure the Locust script to gradually increase the number of users and sustain

load to identify potential bottlenecks. Your script should be able to send 128 images provided to the RESTful API deployed in the Kubernetes cluster.

Ensure the script encodes the images to base64 and embeds them into JSON messages as specified in Section 2 for seamless integration. note: you can reuse part of your client code developed in 2.

7 Experiments and Report - [40 Marks]

Your next objective is to test your system for the maximum load your service can handle under a different number of resources (pods) in your cluster. When the system is up and running, you will run experiments with various *number of pods* (available resources) in the cluster.

You need to conduct two sets of experiments: one where the Locust client runs locally on the master node of Kubernetes, and another where it runs on a VM instance in your pt-project in Nectar pt-xxxxxx project or Azure. The number of pods must be scaled to 1, 2, 3, and 4. Considering the limited CPU and Memory allocated to each pod (CPU request and limit: 0.5, memory request and limit: 512MiB), increasing the number of pods enhances resource accessibility.

Your goal is to determine the maximum number of concurrent users the system can handle before experiencing failures. To achieve this, vary the number of concurrent users in the Locust client to analyze the impact of increased load on the deployed service. You can set the spawn rate to a reasonable value to gradually increase the number of users for various pod configurations. For each trial, continuously send images to the server in a loop until the response time stabilizes and the success rate remains 100%

The response time of a service is the duration between when an end-user makes a request and when a response is sent back. This data is automatically collected by Locust. When the first unsuccessful request occurs, note the maximum number of concurrent users, decrease it by one, and record this number. Then rerun the experiment with the recorded number of concurrent users and a spawn rate of 1 user/second to ensure a 100% success rate.

Finally, report your results along with the average response time in table format, as shown below:

Table 1: Experiment Results

# of Pods	Nectar/Azure		Master	
	Max Users	Avg. Response Time (ms)	Max Users	Avg. Response Time (ms)
1				
2				
4				
8				

Ensure to run each experiment multiple times to verify the correctness of your experiment and the consistency of average response time values across various experiments. This is because network traffic and some other environmental aspects might affect your experiments.

In your report, discuss this table and justify your observations. To automate your experimentation and collect data points, you can write a script that automatically varies the parameters for the experiments and collects data points. Your report should include at least two plots (one for Nectar/Azure, the other for master node). The plots should show the correlation between the number of users and average response time for each number of pods in your experiments. Your plots should have proper title, unit and legends.

Your report must be a maximum of 1500 words excluding your table and references. You need to include the following in your report:

- The table is explained above.
- Explanation of results and observations in your experiments (1000 words).

- Select three challenges of your choice from the list of distributed systems challenges discussed in the first-week seminar, give a practical example from your project that illustrates that challenge and how it is addressed in your system (500 words).

Use 12pt Times font, single column, 1-inch margin all around. Put your **full name**, your **tutor name**, and **student number** at the top of your report.

8 Video Recording

You should submit a video recording and demonstrate your assignment. You should cover the following items in your Video Submission for this assignment:

- Web Service - (approx 2 minutes) Open the source code of your application and briefly explain your program's methodology and overall architecture. Put emphasis on how web service is created, and how JSON messages are created.
- Dockerfile - (approx 1 minute) Briefly explain your approach for containerising the application. Show your Dockerfile, and explain it briefly.
- Kubernetes Cluster and Kubernetes Service - (approx 4 minutes)
 1. Briefly discuss how you installed Docker and Kubernetes and mention which version of these tools are being used. Also, mention which networking module of Kubernetes is used in your setup and why.
 2. List your cluster nodes (`kubectl get nodes, using -o wide`) and explain cluster-info.
 3. Show your deployment YAML file and briefly explain it.
 4. Show your service configuration file and briefly explain it.
 5. Explain and show how your docker image is built and loaded in your Kubernetes cluster.
 6. Show your VMs in the OCI dashboard with your username visible.
 7. Show the public IP address of the controller node and its security group. If you have VCN and subnets you can discuss them as well. Explain why you have configured your security groups and port(s).
 8. For the 4 pods configuration, show that your deployment is working by listing your pods. Then show your service is working and can be reached from outside your controller VM by running the client code on your local computer.
 9. Finally, show the log for pods to demonstrate load balancing is working as expected.
- Locust script - (approx 1 minute) Explain your Locust client and show a quick demo.
- Experiments - There is NO need for any discussion regarding this part in the video.

Caution: Please note that if you do not cover the items requested above in your video you will lose marks even if your code and configurations work properly.

Caution: Your video should be no longer than **8 minutes**. Please note that any content exceeding this duration will result in penalties. Also, kindly refrain from adjusting the recording speed of your video to 1.5x or 2x. The examiners may penalise you if they are unable to follow your talk at a normal pace or understand the content of your presentation.

Recommendation: To ensure that you do not miss any important points in your video recording and stay on track with time, we recommend preparing a script for yourself beforehand. During the recording session, it can be helpful to refer to your script and read through it as needed. You should also prepare all the commands you need to copy and paste before recording.

9 Technical aspects

- Keep your setup up and running during the marking period, as we may access and test your service. Do not remove anything before the teaching team's announcement. Make sure you provide the URL of your service endpoint in the readme.md.
- You can use any programming language. Note that the majority of this project description is written based on Python.
- Make sure you install all the required packages wherever needed depending on the model you use. For example, python, opencv-python, FASTAPI, NumPy and etc.
- When you are running experiments, do not use your bandwidth for other network activities, as it might affect your results.
- Since failure is probable in cloud environments, make sure you take regular backups of your work and snapshots of VMs.
- Make sure your Kubernetes service properly distributes tasks between pods (check logs).
- Make sure you limit the CPU and memory for each pod (0.5 and 512MiB).
- It's important to ensure that your cluster is functioning correctly after each experiment and if redeployment might be necessary in some cases.

10 Submission

You need to submit **four files** via Moodle:

- A report in PDF format as requested (filename: `firstname_lastname_studentid.pdf`).
- A PDF file contains Generative AI prompts you have used and their output.
- A .ZIP file (NOT .RAR or other formats) containing the following:
 1. Your Dockerfile.
 2. Your web service source code.
 3. Your Kubernetes deployment and service configurations (YAML files).
 4. Your Locust Client script.
 5. Any script that automates running experiments if you have one.
- A readme.md markdown file with:
 1. The **URL to a 8-minute video** demonstrating your system. You can use Google Drive, Panopto, or YouTube, e.g., <https://www.youtube.com/watch?v=8frmlor4gTY&t=7s>.
 2. The **URL to your web service endpoints**, e.g, http://118.138.43.2:5000/api/pose_estimation and http://118.138.43.2:5000/api/pose_estimation_annotation

(note: incorrect file naming, format or missing files/URLs may incur a penalty up to 10%.)

Please ensure that your video is private and can be accessed by the teaching team (all tutors and the lecturer). If you would like to inform us about anything else, you can use this ReadMe file. (warning: publicly sharing your assignment video online may assist others to plagiarise and it is considered as an academic misconduct. Severe consequences may apply, please don't take risk.)

11 Generative AI Statement

As per the University's policy on the guidelines and practices pertaining to the usage of Generative AI: AI Generative AI tools may be used OPENLY within this assessment.

Where used, AI must be used responsibly, clearly documented and appropriately acknowledged (see Learn HQ).

Any work submitted for a mark must:

1. Represent a sincere demonstration of your human efforts, skills and subject knowledge that you will be accountable for.
2. Adhere to the guidelines for AI use set for the assessment task.
3. Reflect the University's commitment to academic integrity and ethical behaviour.
4. Inappropriate AI use and/or AI use without acknowledgment will be considered a breach of academic integrity.

The teaching team encourages students to apply their own critical thinking and reasoning skills when working on the assessments with an assistant from GenAI.

Generative AI tools may produce inaccurate content, which could negatively impact student comprehension of big data topics.

All the prompts that you have used and their output must be attached in your submission as a reference.